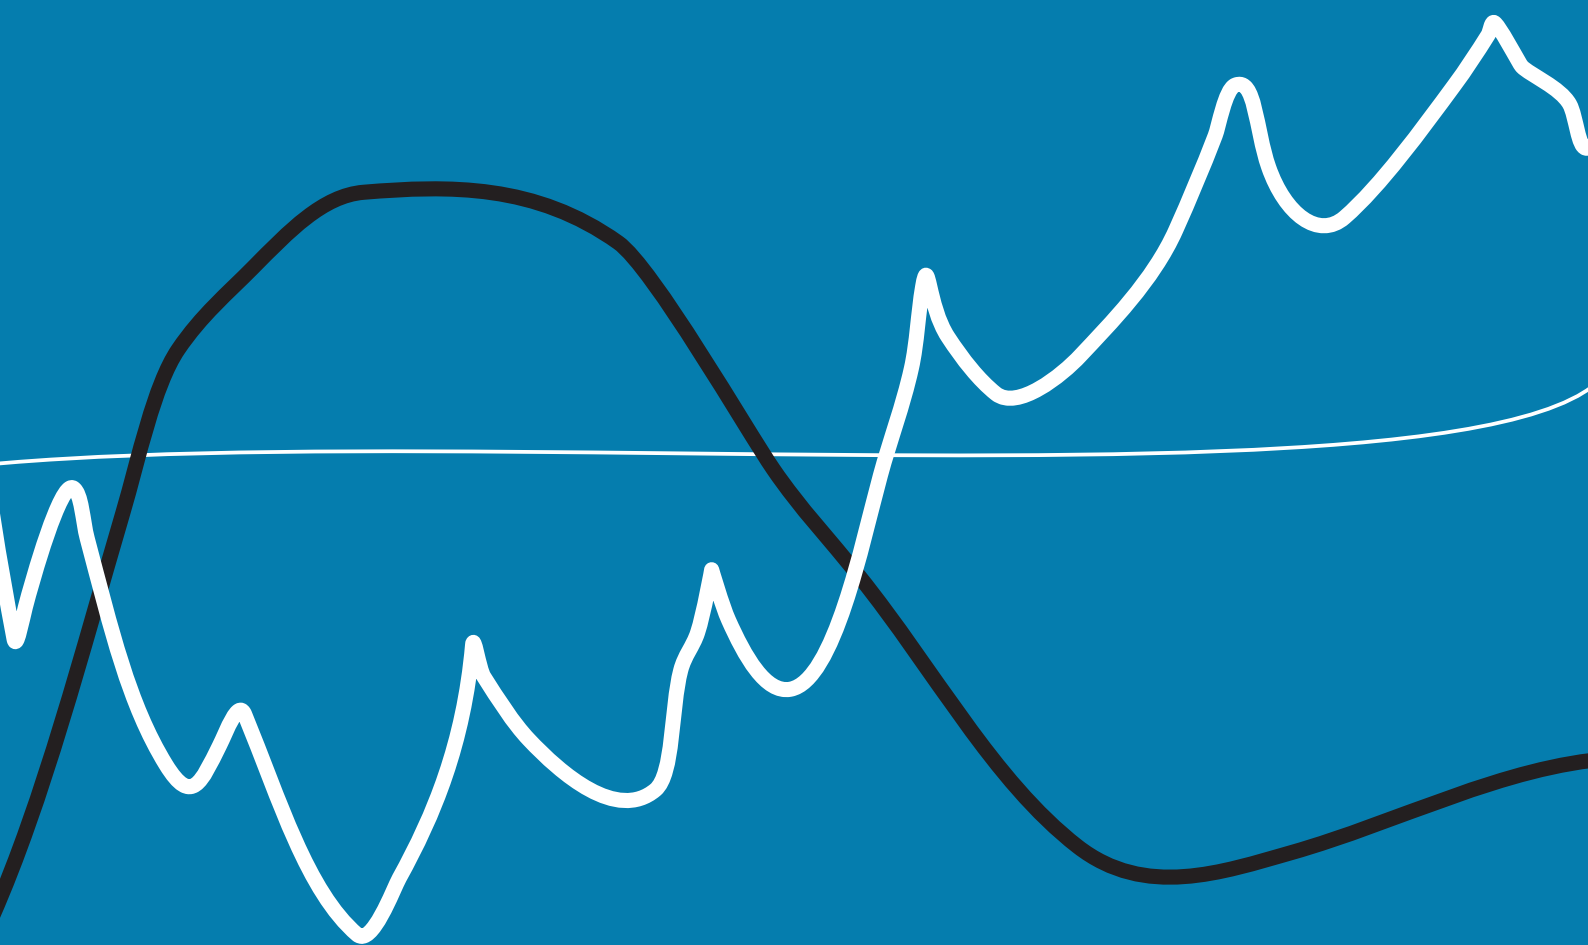


Investigation of Deep Probabilistic Surrogate-models in Bayesian Optimization

Master Thesis



Investigation of Deep Probabilistic Surrogate-models in Bayesian Optimization

Master Thesis
June, 2022

By
Simon Saugmandsgaard Kruse

Copyright: Reproduction of this publication in whole or in part must include the customary bibliographic citation, including author attribution, report title, etc.

Cover photo: Vibeke Hempler, 2012

Published by: DTU, Section of Cognitive Systems, Building 321, 2800 Kgs. Lyngby Denmark
www.github.com/SimonKruse0/master-thesis

ISSN: [0000-0000] (electronic version)

ISBN: [000-00-0000-000-0] (electronic version)

ISSN: [0000-0000] (printed version)

ISBN: [000-00-0000-000-0] (printed version)

Contents

1	Introduction	1
1.1	notation	1
1.2	related work	2
2	Optimization methodology	3
2.1	When to use Bayesian optimization	4
2.2	Observation model	5
3	Bayesian Optimization	7
3.1	Acquisition function	7
3.2	surrogate model	9
4	Gaussian process surrogate	13
4.1	Gaussian Process Regression	13
5	Bayesian neural network surrogate	15
5.1	Bayesian Neural Networks	15
6	Mixture regression surrogate	19
6.1	Gaussian mixture regression	19
6.2	Mixture regression in a Bayesian setting	20
6.3	Sum product networks	20
6.4	SPN	24
6.5	Expectation-maximization for mixture models	27
6.6	Gaussian Mixture Regression	30
7	Results	33
7.1	Mixture regression on simple functions	33
	Bibliography	35
A	Implementation	37

1 Introduction

Optimization plays an important part in our everyday life, science development, and product design. What different transportation should you choose to get fast from A to B, what songs should land on your playlist, and what is the optimal bridge construction. Mathematical optimization problems are all problems in the form,

$$\min_{x \in \mathcal{X}} f(x)$$

where $f : \mathcal{X} \rightarrow \mathcal{R}$ is a functional. I.e if it is possible to set up an objective function. e.g. what is the cost of the bridge given a specific design, or how pleasant you think some music, and some constraints such that you keep in the domain of interest. Evaluation of the objective function can be cheap e.g. if it just requires summing and multiplying numbers or highly expensive if it involves human rating or large simulation and physical experiments. Bayesian optimization is a preferred framework for optimization of the expensive objective functions. And is also referred to as *sample efficient* optimization.

Bayesian optimization is a probabilistic surrogate-based optimization: Assuming some samples from a highly expensive objective a cheap (surrogate) function is used to fit the samples. The next sample is found by minimizing the surrogate and the process is repeated. Bayesian optimization seeks to enhance this procedure with probability theory, where the surrogate function becomes a probabilistic regression model. The most common choice is a Gaussian Process, as it encapsulates the uncertainty very well, but also because its inference procedure (computing answers to probability queries like $p(y|x)$) is exact.

Even though GP has proven good for many cases, there will be problems where the assumptions do not hold. The assumptions of a GP are essential that the objective function can be described as a GP. The nature of a GP is highly dependent on the choice of its kernel and the parameters chosen in that kernel. Another reason for the popularity of the GP is the closed-form inference and giving a closed-form of the expected improvement acquisition function.

The PhD thesis "Sample-efficient Optimization Using Neural Networks" from 2020 [1] showcases empirically that using Bayesian neural networks as surrogate models performed better, or at least comparable to GPs on a wide number of problems. The performance difference was more clear for high-dimensional problems.

This master thesis project will investigate surrogate models alternative to Gaussian processes in Bayesian optimization. Firstly by examining what types of problems a GP surrogate is not a good choice for and where Bayesian neural nets (BNN) surrogates can have an advantage (inspiration found in this 2020 thesis [1]). Secondly by looking at sum-product networks (SPN) as novel surrogate models. An SPN is - similarly to a BNN - a deep probabilistic model and still expressive but with tractable inference, which potentially could lead to advantages over BNNs.

<Make a figure of how the parts are all connected>

1.1 notation

Throughout this thesis we will be using Bayesian notation, i.e. $p(x) := P(X = x)$ is the probability density function of the random variable X evaluated in x . and $p(y|x) := P(Y = y|X = x)$ or $p(y|x) := P(Y|X = x)$.

And writing $p(y^2|x)$ means $P(Y^2 = y^2|X = x)$ and **not** $P(Y = y^2|X = x)$

1.2 related work

This thesis is <BAHAMIANN>

<DNGO> Focus on the inference time of GP scales cubic, which is not appropriate for parallel BayesOpt. Experiments on 6dim Hartmann function.

<Arayns paper.> Conclusions..!?

Already developed alternative surrogate models has been found in the litterateur. The last presented surrogate model, SPN, is to our knowledge not in any published work. More models might still be added to the list.

DNGO

Deep Networks for Global Optimization (DNGO) is presented in the paper 2015 paper "Scalable Bayesian Optimization Using Deep Neural Networks"[2]. The surrogate model is a neural network, where only the last layer is probabilistic, this leads to Bayesian regression and very fast inference.

BOHAMIANN

Bayesian Optimization with **H**amiltonian Monte Carlo Artificial Neural Networks (BOHAMI-ANN) is presented in the 2016 paper "Bayesian Optimization with Robust Bayesian Neural Networks"[3]. This is a fully Bayesian Neural Network trained using adaptive Hamiltonian MCMC.

2 Optimization methodology

Given a cost/objective function $f : \mathcal{X} \rightarrow \mathbb{R}$, where the domain \mathcal{X} could be a subset of \mathbb{R}^n , optimization is a methodology which seeks to find an optimal point, x^* , and value $f^* = f(x^*)$, given as

$$x^* \in \arg \min_{x \in \mathcal{X}} f(x) \quad f^* = \min_{x \in \mathcal{X}} f(x) = f(x^*). \quad (2.1)$$

Note that the above formulation is a minimization problem, which is equivalent to a maximization problem maximizing $-f(\cdot)$. Throughout this thesis, we choose to only work with a minimization problem. Solving this problem is intractable except for rare cases e.g. if f is convex and analytically directly solvable or the domain of f is very limited. Hm

Example: Direct solution method

The unconstrained linear least squares,

$$\min_{x \in \mathbb{R}^n} f(x) := \|Ax - b\|_2^2$$

where $A \in \mathbb{R}^{m \times n}$ and $b \in \mathbb{R}^m$, is a convex problem, i.e. finding x^* such that $\nabla f(x^*) = 0$ is equivalent to finding the solution to the problem. Assuming $A^T A$ is invertible, linear least squares can be solved directly by the normal equations,

$$\nabla f(x) = 2A^T Ax + 2b^T A = 0 \quad \Leftrightarrow \quad x^* = (A^T A)^{-1} A^T b$$

Most optimization problems are non-convex with multiple local minima. And even if the gradient is given analytically the solution is found among a potentially infinitely large set of stationary points ($\nabla f(x) = 0$) and boundary points - this might be tedious or impossible. When the problem is not directly solvable, mathematical optimization takes an indirect methodology: Design a sequence of experiments that reveal information of the objective function. This information can hopefully lead us to the solution of (2.1). This general way of sequentially solving is presented in the book Bayesian Optimization by Roman Garnett [4] and presented here as Algorithm 1. W

Algorithm 1 Sequential Optimization [4]

Input: Initial dataset \mathcal{D} ▷ can be empty
while Temination is not reached **do**
 $x \leftarrow \text{policy}(\mathcal{D})$ ▷ select next observation location
 $y \leftarrow \text{observe}(x)$ ▷ observe objective function at chosen location
 $\mathcal{D} \leftarrow \mathcal{D} \cup \{(x, y)\}$ ▷ update dataset
return: \mathcal{D}

Given data points in the optimization landscape¹ a policy selects a point $x \in \mathcal{X}$ where we make our next observation. Policies can be deterministic or probibalitic, e.g. grid search and random search are examples of policies. The next observation provides us a y value, which combined with x is included to the available data \mathcal{D} . Finally, a stopping criterion decides whether to continue or terminate.

¹”Optimization landscape” defined as the joint set of points in the domain and the objective function evaluated in the points, i.e. $\{(x, f(x)) \in \mathcal{X} \times \mathbb{R} | x \in \mathcal{X}\}$

<example of random search is half as efficient as BO>

Example: Grid search

In grid search values along each dimension in \mathcal{X} is selected and combined with each other, which thereby defines a grid in the parameter space. All points are ordered and systematically selected. In the content of algorithm 1 we define the grid search policy as

$$\text{policy}_{GS}(\mathcal{D}) = x_{|\mathcal{D}|+1}$$

assuming x_1, x_2, \dots, x_n are the ordered grid points and the size of the obtained data is $|\mathcal{D}|$.

Example: Random search

In random search a uniform distribution is laid over the domain space \mathcal{X} and a random point is selected from the distribution.

$$\text{policy}_{RS}(\mathcal{D}) = x, \quad x \sim p(\mathcal{X})$$

Note that grid search and random search are policies which completely ignore the available data. This is a shame and we can do better.

Example: Gradient descent

Gradient descent is the most simple gradient-based optimization approach. The gradient of a continuous function points in the most ascending direction from the location where it is evaluated. In the minimization task, (2.1), we can iteratively use the opposite gradient direction, i.e. the most descending direction. This yields the policy:

$$\text{policy}_{GD}(\mathcal{D}) = x_n - \eta \nabla f(x_n)$$

where we for a brief moment modify y to be a vector, since the observation model is given as:

$$\text{observe}_{GD}(x) = [f(x), \nabla f(x)]$$

Example: Surrogate-based optimization

In surrogate-based optimization all available data is fitted by a cheap-to-evaluate approximation to the objective function - this approximation is called a *surrogate* model. Examples of surrogate models could be a neural network or a random forest. The next point is chosen as the point where the surrogate model is minimized.

$$\text{policy}_{sur}(\mathcal{D}) = \min_x \hat{f}(x)$$

where $\hat{f}(x) \approx f(x)$ for x close to the data \mathcal{D} . And we hope the approximation holds for x far away from the data.

2.1 When to use Bayesian optimization

What if $f(x)$ took several days to evaluate. What if $f(x)$ is noisy? what if discrete points? <more here>




Figure 2.1: Hyper parameter tuning of a model $M(\lambda, \sigma)$, 23 evaluation in grid search vs 23 evaluations using Bayesian optimization

additionally, Bayesian Optimization allow for observation noise,

2.2 Observation model

Many optimization algorithms assume *exact* evaluations of the objective function. However, this assumption is often wrong especially for objective functions with real-life experiments, imperfect simulations, human interaction where measurement noise is a well known. The observation model is typically noisy and described as

$$y = f(x) + \epsilon$$

where ϵ is the measurement error, this is typically assumed to be Gaussian with zero mean and a variance σ^2 (which could depend on x in a heterostodatic setting) and implies a Gaussian observation model,

$$p(y|x, f(x), \sigma) = \mathcal{N}(y; f(x), \sigma^2)$$

we can extend this model to deal with noiseless observations as well, simply by setting $\sigma = 0$ and let the model colaps into a Direct delta distribution,

$$p(y|x, f(x)) = \delta(y - f(x))$$

i.e. all probability mass for y is on the value $f(x)$ giving the observation sample $y = f(x)$

3 Bayesian Optimization

Whereas traditional regression workflow is the following: From data, fit model parameters, make predictions using the parameters. The Bayesian framework allows us to skip the dependency of a single set of parameters and instead use all sets of parameters by treating the set of parameters as a random quantity, θ . What is of interest is the predictive posterior distribution,

$$p(y|x, \mathcal{D}) \quad (3.1)$$


Bayesian optimization is essentially two steps: First, a probabilistic surrogate-model is fitted to the available data \mathcal{D} giving the predictive distribution (3.1). Second, the next sample point is chosen according to a so-called acquisition function, which in a sense balance out the well-known concept, exploitation and exploration. Where exploitation will be choosing the next point according to its expected improvement and exploration will be choosing the next point in a region of high uncertainty and thereby help lower the overall uncertainty. We will first look at the acquisition function used in this thesis, which is also the most commonly used: Expected improvement.

3.1 Acquisition function

A popular choice of acquisition function is expected improvement,

$$EI(x) = \mathbb{E}_{p(y|x, \mathcal{D})}[\max(0, y_{\min} - y)]$$

this will only look at the expectation of the values y , which improves the current best value.



In the following derivation we assume the predictive distribution can be approximated by a normal distribution dependent on the point of interest x and the data \mathcal{D} (note for the GP it is in fact not an

approximation),

$$p(y|x, \mathcal{D}) \approx \mathcal{N}(y|\mu(x, \mathcal{D}), \sigma^2(x, \mathcal{D}))$$

where we will change to a less clumpy notation $\mathcal{N}(y|\mu_x, \sigma_x^2) := \mathcal{N}(y|\mu(x, \mathcal{D}), \sigma^2(x, \mathcal{D}))$. This is completely fine since x is fixed (and \mathcal{D} is fixed) when evaluating the expected improvement in a point x . Furthermore, the density of a standard normal distribution is denoted $\phi(\cdot) := \mathcal{N}(\cdot|0, 1)$, and the cumulative density function (CDF) of a standard normal distribution is denoted, $\Phi(\cdot) := \int_{-\infty}^{\cdot} \phi(\epsilon) d\epsilon$. We will now see that the normal approximation of the predictive distribution yields closed form solution to the expected improvement function,

$$\begin{aligned} E_{p(y|x, \mathcal{D})}[\max(0, y_{\min} - y)] &= \int \max(0, y_{\min} - y) p(y|x, \mathcal{D}) dy \\ &\approx \int \max(0, y_{\min} - y) \mathcal{N}(y|\mu_x, \sigma_x^2) dy \\ &= \int_{-\infty}^{y_{\min}} (y_{\min} - y) \frac{1}{\sigma_x} \phi\left(\frac{y - \mu_x}{\sigma_x}\right) dy \\ &= \int_{-\infty}^{\frac{y_{\min} - \mu_x}{\sigma_x}} (y_{\min} - \mu_x - \sigma_x \epsilon) \frac{1}{\sigma_x} \phi(\epsilon) \sigma_x d\epsilon \\ &= \int_{-\infty}^u \sigma_x \cdot (u - \epsilon) \phi(\epsilon) d\epsilon \\ &= \sigma_x \cdot \left(u \cdot \int_{-\infty}^u \phi(\epsilon) d\epsilon + \int_{-\infty}^u (-\epsilon) \phi(\epsilon) d\epsilon \right) \\ &= \sigma_x [u\Phi(u) + \phi(u)] \end{aligned}$$

where $u := \frac{y_{\min} - \mu_x}{\sigma_x}$. To understand the identity $\phi(u) = \int_{-\infty}^u (-\epsilon) \phi(\epsilon) d\epsilon$ used in the last equality, we first see that the antiderivative is $\phi(\epsilon) = \frac{1}{\sqrt{2\pi}} \exp(-\frac{\epsilon^2}{2})$,

$$\begin{aligned} \frac{d}{d\epsilon} \phi(\epsilon) &= \frac{1}{\sqrt{2\pi}} \frac{d}{d\epsilon} \exp(-\frac{\epsilon^2}{2}) \\ &= \frac{1}{\sqrt{2\pi}} \exp(-\frac{\epsilon^2}{2}) (-\epsilon) \\ &= -\epsilon \phi(\epsilon) \end{aligned}$$

and evaluating the rieman integral is equivalent to evaluate the antiderivative in its boundaries, giving the solution,

$$\int_{-\infty}^u (-\epsilon) \phi(\epsilon) d\epsilon = [\phi(\epsilon)]_{-\infty}^u = \phi(u) - 0 = \phi(u)$$

We can also explicitly write the expected improvement as,

$$EI(x) = (y_{\min} - \mu_x) \Phi\left(\frac{y_{\min} - \mu_x}{\sigma_x}\right) + \sigma_x \phi\left(\frac{y_{\min} - \mu_x}{\sigma_x}\right)$$

where the first part can be interpreted as exploitation (favouring points with a large improvement $I(x) := (y_{\min} - \mu_x)$) and the second part can be seen a exploitation (favouring points with high uncertainty.) taking the derivative with respect to $I(x) := (y_{\min} - \mu_x)$ and σ_x , we see that expected improvement is increasing if the improvement grows or if the variance σ_x grows, i.e

$$\frac{\partial EI(x)}{\partial I(x)} = \Phi\left(\frac{y_{\min} - \mu_x}{\sigma_x}\right) > 0, \quad \frac{\partial EI(x)}{\partial \sigma_x} = \phi\left(\frac{y_{\min} - \mu_x}{\sigma_x}\right) > 0$$

<obs mistake in the book!!!?>

$$\begin{aligned}
\mathbb{E}_{y_*|\mathbf{x}_*, D_n}[\max(0, y_{\min} - y_*)] &= ?? \\
\mathbb{E}[\min(0, y_{\min} - y_*)|\mathbf{x}_*, D_n] &= \int_{-\infty}^{\infty} \min(0, y_{\min} - y_*) p(y_*|\mathbf{x}_*, D_n) dy_* \\
&= \int_{-\infty}^{y_{\min}} (y_{\min} - y_*) p(y_*|\mathbf{x}_*, D_n) dy_* \\
&\approx \frac{1}{N} \sum_{\theta \in \Omega} [y_{\min} - f_{\theta}(x)]
\end{aligned}$$

where $\Omega = \{\theta | f_{\theta}(x) < y_{\min}\}$

3.2 surrogate model

As mentioned in the previous sections, the first of two repeated steps in Bayesian optimization is to create a good Bayesian regression model. i.e. finding the probability of prediction for a arbitrary point x given datapoints $\mathcal{D} = \{x_1, y_1, \dots, x_n, y_n\}$,

$$p(y|x, \mathcal{D})$$

The surrogate model of choice in Bayesian optimization is a Gaussian Process, and Bayesian Neural Network. These are discriminative models, however, another approach, which we focus on in this project, is to model y and x jointly in a so-called generative model.

3.2.1 Discriminative model as surrogate model

When talking about a probabilistic surrogate model we are always implicit talking about a discriminative model: A model of the conditional distribution of the observation, y , conditional on x , i.e.,

$$p(y|x)$$

we also refer to this as the predictive distribution. Gaussian processes and Bayesian neural networks are both discriminative models. There is no distribution over the input x , it is just given. This is indeed sufficient for a surrogate model, where we are interested in the predictive distribution.

3.2.2 Using a generative model as surrogate model

Given a generative model over x and y parametrised with θ , we are dealing with the joint distribution

$$p(x, y|\theta)$$

and we are interested in the conditional distribution of y given x ,

$$p(y|x, \theta_{y|x})$$

where we have put subscript on θ in order to jump up a level of abstraction since, in fact there is just a mapping between them $\theta_{y|x} := g(\theta, y, x)$

$$\begin{aligned}
p(y|x, \mathcal{D}) &= \int p(y|x, \theta_{y|x}) p(\theta_{y|x}|\mathcal{D}) d\theta_{y|x} \\
&= p(y|x, \hat{\theta}_{y|x})
\end{aligned}$$

Where the last equation holds as we assume that $p(\theta_{y|x}|\mathcal{D})$ is a delta function i.e. a point estimate with value $\hat{\theta}_{y|x}$. In the case of our Gaussian mixture model, we obtain a point estimate from the EM algorithm for the variance $\Sigma_{y|k}$, mean value $\mu_{y|k}$ and proportion $\pi_{y|k}$ for each component $k = 1, 2, \dots, K$

$$\hat{\theta}_{y|x} = (\hat{\Sigma}_{y|k}, \hat{\mu}_{y|k}, \hat{\pi}_{y|k})_{k=1}^K$$

However, we are not satisfied with the variance estimate for the regression, as it is way too small for areas with no observed data. It is therefore necessary to manipulate the variance estimate according to that observation. We multiply the variance obtained using expectation-maximization on the joint distribution with the inverse of the probability of the data x , and control that the scaling factor is not going wild!

$$\hat{\Sigma}_{y|k} = \Sigma_{y|k}^{GMM} \frac{1}{\max(p(x), 0.01)}$$

In a way, this is manipulation in a Bayesian spirit, as we let prior and subjective knowledge influence the variance prediction.

3.2.3 Inclusion of prior distribution

The conditional distribution $p(y|x)$ obtained from using a Gaussian mixture model is without a prior, and the predictive distribution yields a way too small uncertainty estimate in regions where there are no observations. This is because the range of the Gaussian distribution continues throughout the entire space and even though the joint distribution is very small in the region, the conditional distribution is normalized using the (very small) $p(x)$. We, therefore, want to introduce a prior, which can rule over the conditional in regions where the probability is too small. The first approach is simply to define a new conditional distribution,

$$\hat{p}(y|x) \propto S(x) \cdot p(y|x) + p_{prior}(y) \quad (3.2)$$

where the scaling function can be interpreted as the probability of the input x is in the region ball $x \in B_{\frac{1}{2}\Delta}(x)$, and the probability density is not just $p(x)$ but the probability of the average mixture component, which is $K \cdot p(x)$, assuming the mixture is dominated by a single component at the position x . We define the scaling as

$$S(x) := K \cdot p(x) \cdot \Delta$$

Next we need to normalize the manipulated predictive distribution in order to make it a proper probability density, i.e., making sure it integrates to 1,

$$\int S(x) \cdot p(y|x) + p_{prior}(y) dy = S(x) + 1$$

so that we can easily update (3.2) with an equality sign by defining the manipulated predictive distribution as,

$$\hat{p}(y|x) = \frac{S(x) \cdot p(y|x) + p_{prior}(y)}{S(x) + 1}.$$

In figure 3.1 the manipulated predictive distribution is illustrated under the influence of different scalings. Note that if $S(x) = 0$ the manipulated distribution becomes the prior, while in the limit $S(x) \rightarrow \infty$ the manipulated distribution becomes the original predictive distribution, $\hat{p}(y|x) = p(y|x)$.

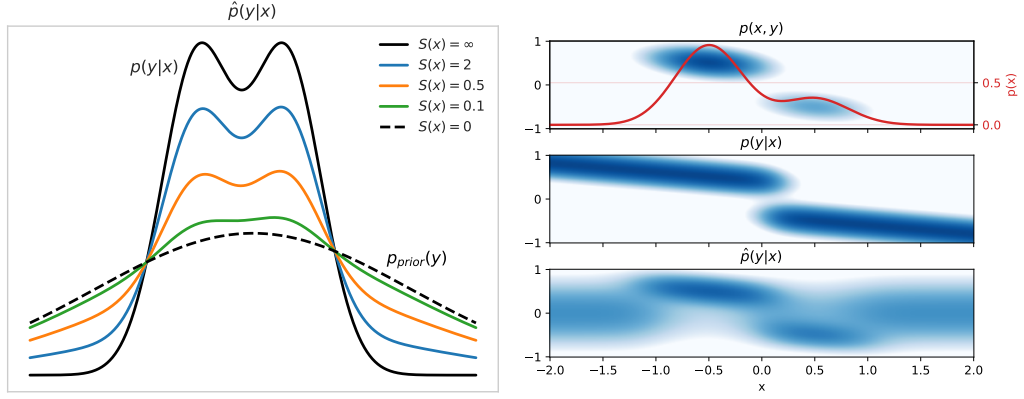


Figure 3.1: Left: Illustration of how the predictive distribution is manipulated according to the scaling function $S(x) := p(x) \cdot N$. Right: Illustration of why it makes sense to manipulate the predictive distribution $p(y|x)$, if there is a small amount of input data at a region, then the predictive distribution should collapse into the uncertain prior

Manipulated predictive mean

What is the mean prediction of the manipulated predictive distribution?

$$\begin{aligned}
 E_{p_{bayes}(y|x)}[y] &= \int y \cdot \frac{f(n, p(x)) \cdot p(y|x) + \lambda p_{prior}(y)}{Z} dy \\
 &= \left(p(x) \cdot N \cdot E_{p(y|x)}[y] + \lambda \cdot E_{p_{prior}(y)}[y] \right) \frac{1}{Z} \\
 &= \frac{p(x) \cdot N \cdot E_{p(y|x)}[y]}{p(x) \cdot N + \lambda}
 \end{aligned}$$

Manipulated predictive variance

What is the variance of the manipulated predictive distribution We first calculate the second moment,

$$\begin{aligned}
 E_{p_{bayes}(y|x)}[y^2] &= \int y^2 \cdot \frac{p(x) \cdot N \cdot p(y|x) + \lambda p_{prior}(y)}{Z} dy \\
 &= (p(x) \cdot N \cdot E_{p(y|x)}[y^2] + \lambda E_{p_{prior}(y)}[y^2]) \frac{1}{Z} \\
 &= \frac{p(x) \cdot N \cdot (Var_{p(y|x)}[y] + E_{p(y|x)}[y]^2) + \lambda Var_{p_{prior}(y)}[y]}{p(x) \cdot N + \lambda}
 \end{aligned}$$

So the variance is calculated as,

$$V_{p_{bayes}(y|x)}[y] = E_{p_{bayes}(y|x)}[y^2] - E_{p(x,y)}[y]^2$$

implementation

It is not necessary to calculate the conditional distribution, since, using Bayes rule

$$p(y|x) = \frac{p(x, y)}{p(x)}$$

so gives

$$p_{bayes}(y|x) = \frac{p(x) \cdot N \cdot p(y|x) + \lambda p_{prior}(y)}{Z} = \frac{N \cdot p(x, y) + \lambda p_{prior}(y)}{Z}$$

$$E_{p(x,y)}[y] = \int_y \int_x yp(x, y) dx dy = \int_y yp(y) dy = E_{p(y)}[y]$$

4 Gaussian process surrogate

4.1 Gaussian Process Regression

We now show how the predictive distribution is calculated exact for Gaussian Processes, i.e.

$$p(y|x, \mathcal{D}) = \int \mathcal{N}(y|f(x), \sigma^2) p(f(x)|\mathcal{D}) df(x) \quad (4.1)$$

we will soon see that $p(f(x)|\mathcal{D}) = \mathcal{N}(f(x)|\dots)$ and thereby that we have a marginal Gaussian distribution for $f(x)$ and a conditional Gaussian distribution of y given $f(x)$, giving us the marginalized distribution, $p(y|x, \mathcal{D})$, using formulas (4.2).

Background: Trick with normal distributions [from Bishops book?]

Given a marginal Gaussian distribution of x and a conditional Gaussian distribution of y given x of the form,

$$\begin{aligned} p(x) &= \mathcal{N}(x|\mu, \Lambda^{-1}) \\ p(y|x) &= \mathcal{N}(y|Ax + b, L^{-1}) \end{aligned}$$

then the marginal distribution of y and the conditional distribution of x given y have the form,

$$p(y) = \mathcal{N}(y|A\mu + b, L^{-1} + A\Lambda^{-1}A^T) \quad (4.2)$$

$$p(x|y) = \mathcal{N}(x|\Gamma\mu + \Gamma[A^T L(y - b)], \Gamma) \quad (4.3)$$

$$\Gamma := (\Lambda + A^T L A)^{-1} \quad (4.4)$$

Posterior function

Recall we assume $\mathbf{f} = (f(\mathbf{x}_1), \dots, f(\mathbf{x}_n))$ is the parameters in our model, therefore we call $p(\mathbf{f}|\mathcal{D})$ the posterior distribution. However, what is of real interest is the function values on unobserved locations, thereby we extend \mathbf{f} to be a function, i.e. an infinitely dimensional vector. We call this quantity *the posterior function*

$$p(f(\cdot)|\mathcal{D}) = \int p(f(\cdot)|\mathbf{x}, \mathbf{f}) p(\mathbf{f}|\mathcal{D}) d\mathbf{f}. \quad (4.5)$$

Prior we assume that the function takes values according to

$$p(\mathbf{f}|\mathbf{x}) = \mathcal{N}(\mathbf{f}|\mathbf{0}, c(\mathbf{x}, \mathbf{x}))$$

where the covariance is defined at kernel evaluation for each pair of \mathbf{x} , where $c(\cdot, \cdot)$ is a covariance function, chosen to be the Matérn covariance function,

$$c(\mathbf{x}, \mathbf{x}) = \begin{bmatrix} c(x_1, x_1) & \dots & c(x_1, x_n) \\ \vdots & \ddots & \vdots \\ c(x_n, x_1) & \dots & c(x_n, x_n) \end{bmatrix} \quad c(x, y) := \text{Matern}(x, y) \dots$$

Appendix <...> tries to give the intuition why this makes sense.

We now calculate the first term in the integral (4.5), $p(f(\cdot)|\mathbf{x}, \mathbf{f})$ using that we have the joint prior distribution,

$$p(f(\cdot), \mathbf{f}|\mathbf{x}) = \mathcal{N}\left(\begin{bmatrix} f(\cdot) \\ \mathbf{f} \end{bmatrix} \middle| \begin{bmatrix} 0 \\ \mathbf{0} \end{bmatrix}, \begin{bmatrix} c(\cdot, \cdot) & c(\cdot, \mathbf{x}) \\ c(\mathbf{x}, \cdot) & c(\mathbf{x}, \mathbf{x}) \end{bmatrix}\right) \quad (4.6)$$

And the conditonal of a joint Gaussian is given using <ref>

$$p(f(\cdot)|\mathbf{x}, \mathbf{f}) = \mathcal{N}(f(\cdot)|c(\cdot, \cdot)^{-1}c(\cdot, \mathbf{x})\mathbf{f}, c(\cdot, \cdot)^{-1})$$

Next we calculate the last term in the integral (4.5), $p(\mathbf{f}|\mathcal{D})$, i.e. the posterior distribution. Assuming iid data, i.e. $p(y_1, \dots, y_n|x_1, \dots, x_n, \mathbf{f}) = \prod_{i=1}^n p(y_i|x_i, \mathbf{f}_i)$ and that the likelihood is Gaussian with mean \mathbf{f} and variance $\sigma^2 I_n$.

$$\begin{aligned} p(\mathbf{f}|\mathcal{D}) &\propto p(\mathbf{f}|\mathbf{x}) \prod_{i=1}^n p(y_i|x_i, \mathbf{f}_i) \\ &= \mathcal{N}(\mathbf{f}|\mathbf{0}, c(\mathbf{x}, \mathbf{x})) \prod_{i=1}^n \mathcal{N}(y_i|\mathbf{f}_i, \sigma^2) \\ &= \mathcal{N}(\mathbf{f}|\mathbf{0}, c(\mathbf{x}, \mathbf{x})) \mathcal{N}(\mathbf{y}|\mathbf{f}, \sigma^2 I_n) \end{aligned}$$

now from <ref> we have that the posterior is the following Gaussian:

$$p(\mathbf{f}|\mathcal{D}) = \mathcal{N}(\mathbf{f}|M^{-1}\sigma^{-2}\mathbf{y}, M^{-1}) \quad M := c(\mathbf{x}, \mathbf{x})^{-1} + \sigma^{-2}I_n$$

Now we found that both term in the integral (4.5), and they are related such that it is possible to use (4.2) for arriving at (we define $A := c(\cdot, \cdot)^{-1}c(\cdot, \mathbf{x})$),

$$p(f(\cdot)|\mathcal{D}) = \mathcal{N}(f(\cdot)|AM^{-1}\sigma^{-2}\mathbf{y}, c(\cdot, \cdot)^{-1} + AM^{-1}A^T)$$

Finally we found that both terms in the integral (4.1) also is related in a similar way, and we use (4.2), again to arrive at the predictive distribtuion,

$$p(y_*|x_*, \mathcal{D}) = \mathcal{N}(y_*|AM^{-1}\sigma^{-2}\mathbf{y}, c(x_*, x_*)^{-1} + AM^{-1}A^T + \sigma^2)$$

Some questions about a naive approach..!

Learning - Emperical bayes inference

Anther inference which is done is then optimizing the hyper parameters using emperical bayes i.e. the variance and length scale for the kernel. Here we optimize the marginalized likelihood function

$$p(y_1, \dots, y_n|x_1, \dots, x_n, \theta) = -\frac{1}{2}[(y - \mu)^T(\Sigma + N)^{-1}(y - \mu) + \log |\Sigma + N| + n \log 2\pi]$$

<and how to get to there?>

Model assessment becomes trivial in light of the model posterior if we simply establish preferences over models according to their posterior probability. When using the uniform model prior (4.6) the model posterior is proportional to the marginal likelihood alone, which can be then used directly for model assessment. ??! Youtube - good video! Go through that!

5 Bayesian neural network surrogate

5.1 Bayesian Neural Networks

Bohamiann and numpyro-BNN are examples of probabilistic models with intractable inference. The predictive density is given as,

$$\begin{aligned} p(y_*|x_*, \mathcal{D}) &= \int p(y_*|x, \theta) p(\theta|\mathcal{D}) d\theta \\ &\approx \frac{1}{K} \sum_{k=1}^K p(y_*|x, \theta^{(k)}) \end{aligned}$$

where the first integral is intractable as θ can live in a highly dimensional space, second the approximation sign is true, since we can approximate the integral with monte carlo sampling: $\theta^{(k)}$ are iid samples from the posterior distribution, $\theta^{(k)} \sim p(\theta|\mathcal{D})$. We can get samples from the posterior distribution

Background: Monte Carlo approximation

Assuming we have a number of iid samples, $\theta^{(1)}, \dots, \theta^{(K)}$ drawn from the distribution $p(x)$, then the following approximation

$$E[f(x)] \approx \frac{1}{K} \sum_{k=1}^K f(x^{(k)}) =: \Theta_K(f)$$

holds according to the law of large numbers in fact

$$E[f(x)] = \lim_{K \rightarrow \infty} \Theta_K(f)$$

and the central limit theorem, <OBS refere!>

$$p(\hat{\Theta}) \approx \mathcal{N}(\hat{\Theta}|\mu_f, \frac{\sigma_f^2}{K})$$

which ensures that the variance of the unbiased estimator of the expectation decreases with number of samples, K . Left is to sample the *iid* samples from the distribution $p(x)$

Posterior samples

For both models the joint distribution $p(\mathcal{D}, \theta)$ is available, but calculating the posterior distribution requires the marginalized likelihood, $p(\mathcal{D}) = \int_{\theta} p(\mathcal{D}, \theta)$. This integral is often intractable since the space of θ typically is abnalous - so not even numerical approximations of the integral is tractable. From Bayes rule, we have the equality,

$$p(\theta|\mathcal{D}) = \frac{p(\mathcal{D}, \theta)}{p(\mathcal{D})} \propto p(\mathcal{D}, \theta),$$

where the propotional sign is true, since $p(\theta|\mathcal{D})$ is a function of θ . Knowing the $p(\mathcal{D}, \theta)$ joint distribution thereby allow for using Markov chain Monte Carlo for sampling from the posterior distribution.

Background: Markov chain Monte Carlo

We can conveniently use MCMC for sampling from a probability density $p(x)$, with only the knowledge of a proportional/unnormalised density $\hat{p}(x) \geq 0$ i.e

$$\hat{p}(x) = c \cdot p(x) \quad c = \int \hat{p}(x) dx,$$

where $\int \hat{p}(x) dx$ is a possible intractable integral. An ergodic Markov chain/process is constructed, such that its stationary distribution is exactly $p(x)$, but only with the knowledge of $\hat{p}(x)$.

Example: Metropolis-Hasting (MH)

The most simple MCMC method is the Metropolis-Hasting algorithm. At iteration n we have a sample x_n ,

1. Propose \hat{x} from a proposal density $q(x_n, \cdot)$
2. Compute acceptance probability

$$\alpha(x_n, \hat{x}) = \min \left(1, \frac{p(\hat{x})}{p(x_n)} \frac{q(x_n, \hat{x})}{q(\hat{x}, x_n)} \right)$$

3. Set the next sample

$$x_{n+1} = \begin{cases} \hat{x} & \text{with probability } \alpha(x_n, \hat{x}) \\ x_n & \text{with probability } 1 - \alpha(x_n, \hat{x}) \end{cases}$$

note that $\alpha(x_n, \hat{x})$ requires $p(x)$, but since the algorithm only requires the fraction $\frac{p(\hat{x})}{p(x_n)} = \frac{p(\hat{x}) \cdot c}{p(x_n) \cdot c} = \frac{\hat{p}(\hat{x})}{\hat{p}(x_n)}$ we only need \hat{p} .

Proof: Assuming discrete states, the transition probability between the states are given as,

$$p(x \rightarrow y) = \begin{cases} q(x, y) \alpha(x, y) & \text{if } x \neq y \\ q(x, x) + \sum_{z \neq x} q(x, z) (1 - \alpha(x, z)) & \text{if } x = y \end{cases}$$

Now, let us look at the so-called *detailed balance* relation, i.e. that if we are sampling from the stationary density we stay there at the next state. Assume $x \neq y$,

$$\begin{aligned} p(x)p(x \rightarrow y) &= p(x)q(x, y)\alpha(x, y) \\ &= p(x)q(x, y) \min \left(1, \frac{p(\hat{x})}{p(x_n)} \frac{q(x_n, \hat{x})}{q(\hat{x}, x_n)} \right) \\ &= \min(p(x)q(x, y), p(y)q(y, x)) \end{aligned}$$

Observing that the right hand side yields symmetric result in x and y , therefore we obtain,

$$p(x)p(x \rightarrow y) = p(y)p(y \rightarrow x)$$

and summing over x on both sides yields,

$$\sum_x p(x)p(x \rightarrow y) = p(y) \sum_x p(y \rightarrow x) \quad (5.1)$$

$$\implies p(y) = \sum_x p(x)p(x \rightarrow y) \quad (5.2)$$

similar conclusion will be obtained for $x = y$, all in all this reveals that $p(x)$ is in fact invariant for the chain $\{x_1, \dots, x_n\}$ and thereby that MH is a MCMC algorithm.

HM with random walk transition is very simple and it comes with some serious disadvantages: slow convergence speed, might stay in the same region for a long time and produces highly correlated samples. We can do better by replacing the random walk with gradient-guided movements and interpretate the probability landscape as a physical system.

Example: HMC

The golden standard in MCMC is the Hamilton monte carlo, which exploits arguments from classical mechanics around the Hamiltonian equations. This method leads to more efficient sampling as the Hamiltonian interpretation allows the system to take regions with high probability mass into account - this is obtained using gradient of the probability landscape. $\frac{-\partial E(x)}{\partial x}$

We define PDF

$$p(x) = \frac{1}{Z_E} \exp(-E(x)),$$

where $E(x)$ is interpreted as the systems potential energy. Now, an latent vector q is introduced in order to represent the momentum of the system, which gives us the kinetic energy of the system.

$$K(q) = \frac{1}{2} \sum_{i=1}^l q_i^2$$

Giving the Hamilton function and its corresponding distribution

$$H(x, q) = E(x) + K(q)$$

and

$$p(x, q) = \frac{1}{Z_H} \exp(-H(x, q)) \quad (5.3)$$

$$= \frac{1}{Z_E} \exp(-E(x)) \frac{1}{Z_K} \exp(-K(x)) \quad (5.4)$$

$$= p(x)p(q) \quad (5.5)$$

The desired distribution $p(x)$ is found as the marginal of $p(x, q)$

since some intuition is now established around Hamiltonian Monte Carlo, we look a bit on the two versions used in Numpyro-BNN and Bohamiann,

5.1.1 No U-Turn sampling

often the physical simulation in HMC goes forth and back the same path, and we risk getting bad samples. No U-turn (NUTS) sampling avoid this.

5.1.2 Adaptive stochastic HMC

6 Mixture regression surrogate

6.1 Gaussian mixture regression

Taking a convex combination of a set of multivariate Gaussian distributions is a Gaussian mixture model

$$p(z) = \sum_{k=1}^K \pi_k \mathcal{N}(z | \mu_k, \Sigma_k)$$

Defining $z := (x, y)$ we can model our data, as a generative model $p(x, y)$, now, since the conditional of a Gaussian mixture again is Gaussian mixture - i.e. closed form expression, we can exactly calculate $p(y|x) = GMM_{y|x}$

Assuming iid data the likelihood is given as

$$p(\mathcal{D} | \mu_1, \dots, \mu_K, \Sigma_1, \dots, \Sigma_K, \pi_1, \dots, \pi_K) = \prod_{i=1}^n \sum_{k=1}^K \pi_k \mathcal{N}(z_i | \mu_k, \Sigma_k)$$

And the log likelihood,

$$\log p(\mathcal{D} | \cdot) = \sum_{i=1}^n \log \sum_{k=1}^K \pi_k \mathcal{N}(z_i | \mu_k, \Sigma_k)$$

Dealing with wrong variance

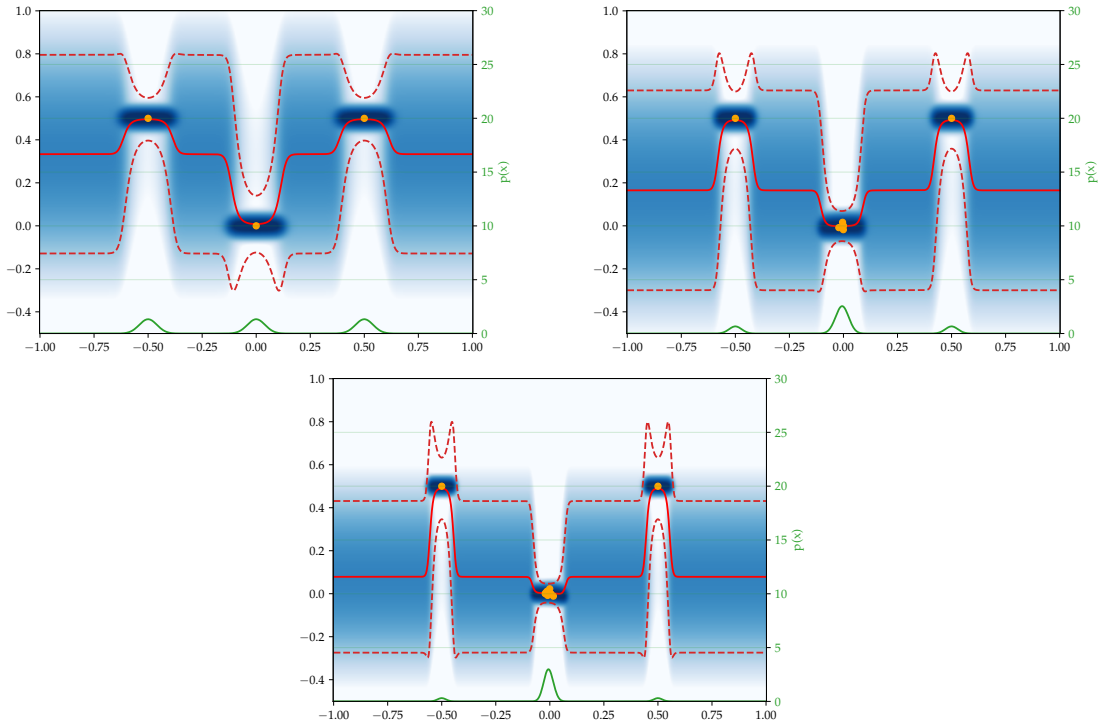


Figure 6.1: EI improvement might in this case be stuck at the point 0, since the conditional distribution is normalized by $p(x)$ and therefore not influenced by the amount of data

<variance manipulation> <Expected improvement manipulation>

6.2 Mixture regression in a Bayesian setting

As seen in examples. The uncertainty of conditional distribution is way too certain in areas with no data points, therefore we need to enhance the model with some bayesian flavour.

$$p(y|x, \mathcal{D}) = p(y|x, Z)p(Z|x)$$

Expetation-maximization algorithm

A way to find local maxima in the likelihood function is using the EM algorithm.

If we define a latent/hidдем random variabel $Z_i \in \{1, \dots, K\}$ for each data point, then the likelihood function becomes,

$$L(\theta|\mathcal{D}, Z) = \prod_{i=1}^n \sum_{k=1}^K 1(Z_i = k) \pi_k \mathcal{N}(z_i | \mu_k, \Sigma_k)$$

Now the expectation wrt. the current value $p(Z|\mathcal{D}, \theta^k)$ is given as

$$Q(\theta|\theta^k) = \mathcal{E}_{p(Z|\mathcal{D}, \theta^k)} = L(\theta|\mathcal{D}, Z)$$

And then update the next parameter estimate with

$$\theta^{k+1} = \arg \max_{\theta} Q(\theta|\theta^k)$$

This is repeated untill convergence.

6.3 Sum product networks

We will for a large extend just see SPN as a large mixture model. This is a totally valid observation.

Mixture model

Definition 1 A sub-network \bar{S}_z of S is an SPN, which includes the root S and then includes nodes according to the following recursive scheme:

Collection of sub-network S_z of S

function Process(node i , S_z)

if $i \in \mathcal{Leaf}(S)$ **then**

return:

if $i \in \mathcal{Sum}(S)$ **then**

$S_z = S_z.add(j \in ch(i))$

 ▷ include one child of node i

return: Process(j , S_z)

if $i \in \mathcal{Prod}(S)$ **then**

$S_z = S_z \cup \{j | j \in ch(i)\}$

 ▷ include all children of node i

for $j \in ch(i)$ **do**

return: Process(j , S_z)

return: S_z

$S_z = \text{Process}(\text{root}, \emptyset)$

So we see that at each sum node the number of different sub-networks multiplies with the number of children for that sum node. And thereby, the total number of sub-networks is

$$Z = \prod_{i \in \mathcal{Sum}(S)} |ch(i)|$$

i.e. an exponential large amount of sub-networks. This is the amount of mixture components implicitly defined in an SPN. Denote the set of edges in the sub-network $\mathcal{E}(S_z)$. Now we define a mixture coefficient, λ_z and component for each S_z as

$$\lambda_z := \prod_{(i,j) \in \mathcal{E}(S_z)} w_{i,j}, \quad p_z(x, y|\theta) := \prod_{i \in \mathcal{L}(S_z)} p_i(x, y)$$

where $p_i(x, y)$ is the leaf distribution at leaf node i parametrised with θ . It can now be proven that the SPN can be interpreted as the following mixture model,

$$p(x, y|w, \theta) = \sum_{z=1}^Z \lambda_z(w) p_z(x, y|\theta)$$

i.e. by the weighted sum of all Z sub-networks. For convenience we define each sum component as $p(z, x, y|w, \theta) := \lambda_z(w) p_z(x, y|\theta)$. Evaluation of $p(x, y|w, \theta)$ will never be done as the sum over Z components, instead there is a proposition.

Proposition 1 Consider a SPN, S , a sum node $q \in \text{Sum}(S)$ and a child $i \in \text{ch}(q)$, then the following relation holds,

$$\sum_{z: (q,i) \in \mathcal{E}(S_z)} \lambda_z(w) p_z(x, y|\theta) = w_{i,q} \frac{\partial S}{\partial v(q)} v(i)$$

Conditional of SPN

We will soon see how it is possible to write the conditional distribution as the mixture,

$$p(y|x) = \sum_{z \in \Sigma(S)} \gamma(x) p_{z_y}(y)$$

where $\Sigma(S)$ is the set of all sub-networks in the SPN, S - **IT IS EXPONENTIALLY LARGE**. And where $p_{z_y}(y)$ is defined through $p_z(x, y)$,

$$\begin{aligned} p_z(x, y) &= \prod_{l \in \text{Leaf}(z_x)} \phi_l(x) \prod_{l \in \text{Leaf}(z_y)} \phi_l(y) \\ &=: p_{z_x}(x) p_{z_y}(y) \end{aligned}$$

where ϕ_l is the density of the l 'th leafs tractable distribution. Recall that we can interpret an SPN as the mixture model,

$$p(x, y) = \sum_{z \in \Sigma(S)} \lambda_z p_z(x, y)$$

where $\lambda_z = \prod_{(q,j) \in \mathcal{E}(z)} w_{q,j}$. First we calculate the marginal density, $p(x)$,

$$\begin{aligned} p(x) &= \int p(x, y) dy \\ &= \int \sum_{z \in \Sigma(S)} \lambda_z p_z(x, y) dy \\ &= \sum_{z \in \Sigma(S)} \lambda_z p_{z_x}(x) \int p_{z_y}(y) dy \\ &= \sum_{z \in \Sigma(S)} \lambda_z p_{z_x}(x) \end{aligned}$$

Now we are ready to calculate the conditional density,

$$\begin{aligned}
p(y|x) &= \frac{p(x, y)}{p(x)} \\
&= \frac{\sum_{z \in \Sigma(S)} \lambda_z p_z(x, y)}{p(x)} \\
&= \sum_{z \in \Sigma(S)} \frac{\lambda_z p_{z_x}(x)}{p(x)} p_{z_y}(y) \\
&= \sum_{z \in \Sigma(S)} \frac{\lambda_z p_{z_x}(x)}{\sum_{z \in \Sigma(S)} \lambda_z p_{z_x}(x)} p_{z_y}(y) \\
&= \sum_{z \in \Sigma(S)} \gamma(x) p_{z_y}(y)
\end{aligned}$$

So we defined $\gamma(x) = \frac{\lambda_z p_{z_x}(x)}{\sum_{z \in \Sigma(S)} \lambda_z p_{z_x}(x)}$ and this is very convinient, as we will see soon is equivalent to the derivative of the log-likelihood of the SPN, which is easily obtained by automatic differentiation.

calculation of $\gamma(x)$

expectation maximization of a mixture model, is given by Bishop.. the responsibility of a datapoint to belong to one mixture component, is given by

$$\gamma(z_{nk}) = \frac{w_k p_j(x_n)}{\sum_i w_i p_i(x_n)}$$

We can prove that the responsibility is equal to the gradient of the log likelihood,

$$L := \sum_n \log \sum_j w_j \exp \psi_j(x_n)$$

where we define $\psi_j(x_n) = \log p_j(x_n)$. Take the gradient

$$\frac{\partial L}{\partial \psi_j(x_n)} = \frac{w_k p_j(x_n)}{\sum_i w_i p_i(x_n)}$$

Note that the gradient easily can be found using autograd.

Mean and variance of $p(y|x)$

The mean of the conditional is just

$$\begin{aligned}
E_{p(y|x)}[y] &= \sum_{z \in \Sigma(S)} \gamma(x) \int y p_{z_y}(y) dy \\
&= \sum_{z \in \Sigma(S)} \gamma(x) \prod_{l \in \text{Leaf}(z_y)} E_{\phi_l}[y]
\end{aligned}$$

and the variance is found using the second moment,

$$\begin{aligned}
E_{p(y|x)}[y^2] &= \sum_{z \in \Sigma(S)} \gamma(x) \int y^2 p_{z_y}(y) dy \\
&= \sum_{z \in \Sigma(S)} \gamma(x) \prod_{l \in \text{Leaf}(z_y)} (Var_{\phi_l}[y] + E_{\phi_l}[y]^2)
\end{aligned}$$

SPN is an exponential large mixture model, with linear inference - unlike GMM. !? Write naive bayesian mixture model as a Sum Product Network

This thesis will only work with RAT spn. Which takes combinations of each dimension. and...

sum nodes play a role of mixtures over their children distribution, similar to a classic mixture model

Product nodes on the other hand, are equivalent to factorizations over independent distributions as they are combining disjoint RVs.

SPNs can also be interpreted as deep feed forward neural network [vergari]. Here, imagine the weights of the sum nodes are parameters, leaf distributions are input neurons, root node is output and all other nodes correspond to hidden neurons




Figure 6.2: SPN





Figure 6.3: Model including wifi information



6.4 SPN

Sum-product networks are generative models, i.e. statistical models of the joint distribution $p(x, y)$. We need, however, a discriminative model for regression, i.e. a model of the conditional distribution $p(y|x)$. SPNs allow for exact inference of the joint distribution and any marginalized distribution. These combined queries is sufficient for the exact predictive posterior.

6.4.1 SPN - prediction

Prior to the inference of the predictive distribution, we assume that the SPN, S , is trained, i.e. trained leaf distributions $p_j(\cdot)$ for all leaf nodes, $j \in \text{Leaf}(S) := \{j \in \mathcal{V}(S) | \text{pa}(j) = \emptyset\}$ and weights $w_{i,j}$ for the connections between every sum nodes $i \in \mathcal{S}$ and its children, $j \in \text{ch}(i)$.

The joint and the marginal distribution are evaluated in the following recursive way

Calculation of $p(x, y)$

Input: Fully trained SPN, with leaf distributions $p_i(\cdot)$ for $i \in \text{Leaf}(S)$ and weights $w_{i,j}$ for $(i, j) \in \{(i, j) | i \in \text{Sum}(S), j \in \text{ch}(i)\}$

function Eval(node i)

if $i \in \text{Leaf}(S)$ **then**

return: $p_i(x, y)$

 ▷ evaluate leaf distributions at point (x, y)

if $i \in \text{Sum}(S)$ **then**

return: $\sum_{j \in \text{ch}(i)} w_{i,j} \text{Eval}(j)$

if $i \in \text{Prod}(S)$ **then**

return: $\prod_{j \in \text{ch}(i)} \text{Eval}(j)$

$p(x) = \text{Eval}(\text{root})$

Calculation of $p(x)$

Input: Fully trained SPN, with leaf distributions $p_i(\cdot)$ for all leaves i and weights w_i .

function Eval(node i)

- if** $i \in \mathcal{Leaf}(S)$ **then**
 - if** node handle x **then**
 - return:** $p_i(x, y)$ \triangleright evaluate leaf distributions at point (x, y)
 - else**
 - return:** 1 \triangleright set node equal 1 at point (x, y)
- if** $i \in \mathcal{Sum}(S)$ **then**
 - return:** $\sum_{j \in ch(i)} w_{i,j} \text{Eval}(j)$
- if** $i \in \mathcal{Prod}(S)$ **then**
 - return:** $\prod_{j \in ch(i)} \text{Eval}(j)$

$p(x) = \text{Eval}(\text{root})$

So after doing two slightly different forward passes through the SPN, $p(x)$ and $p(x, y)$, using Bayes rule, we can combined the two queries into the conditional distribution:

$$p(y|x) = \frac{p(x, y)}{p(x)}$$

The predictive distribution is found with a cost of just $O(E + E + 1) = O(E)$, where E is number of edges/connections in the SPN.

6.4.2 SPN - learning

It is not enough to do predictive inference on a SPN, we also need to fit it on the data. It is possible interpret sum-product network as a large mixture model and therefore use expectation-maximization to train the model. We will introduce that idea now. The Paper [SPN_EM]... defines SPN as a mixture of all sub-networks of an SPN.

Definition 2 A sub-network \bar{S}_z of S is an SPN, which includes the root S and then includes nodes according to the following recursive scheme:

Collection of sub-network S_z of S

function Process(node i , S_z)

- if** $i \in \mathcal{Leaf}(S)$ **then**
 - return:**
- if** $i \in \mathcal{Sum}(S)$ **then**
 - $S_z = S_z.add(j \in ch(i))$ \triangleright include one child of node i
 - return:** Process(j , S_z)
- if** $i \in \mathcal{Prod}(S)$ **then**
 - $S_z = S_z \cup \{j | j \in ch(i)\}$ \triangleright include all children of node i
 - for** $j \in ch(i)$ **do**
 - return:** Process(j , S_z)

return: S_z

$S_z = \text{Process}(\text{root}, \emptyset)$

So we see that at each sum node the number of different sub-networks multiplies with the number of children for that sum node. And thereby, the total number of sub-networks is

$$Z = \prod_{i \in \text{Sum}(S)} |ch(i)|$$

i.e. an exponential large amount of sub-networks. This is the amount of mixture components implicitly defined in an SPN. Denote the set of edges in the sub-network $\mathcal{E}(S_z)$. Now we define a mixture coefficient, λ_z and component for each S_z as

$$\lambda_z := \prod_{(i,j) \in \mathcal{E}(S_z)} w_{i,j}, \quad p_z(x, y|\theta) := \prod_{i \in \mathcal{L}(S_z)} p_i(x, y)$$

where $p_i(x, y)$ is the leaf distribution at leaf node i parametrised with θ . It can now be proven that the SPN can be interpreted as the following mixture model,

$$p(x, y|w, \theta) = \sum_{z=1}^Z \lambda_z(w) p_z(x, y|\theta)$$

i.e. by the weighted sum of all Z sub-networks. For convenience we define each sum component as $p(z, x, y|w, \theta) := \lambda_z(w) p_z(x, y|\theta)$. Evaluation of $p(x, y|w, \theta)$ will never be done as the sum over Z components, instead there is a proposition.

Proposition 2 Consider a SPN, S , a sum node $q \in \text{Sum}(S)$ and a child $i \in ch(q)$, then the following relation holds,

$$\sum_{z: (q,i) \in \mathcal{E}(S_z)} \lambda_z(w) p_z(x, y|\theta) = w_{i,q} \frac{\partial S}{\partial v(q)} v(i)$$

6.5 Expectation-maximization for mixture models

Mixture models can be seen as probabilistic graphical models, <fig> there one mixture component is picked according to the realization of a categorical variable \mathbf{Z} with parameters according the the mixture weights, i.e we can reformulate,

$$p(x) = \sum_{k=1}^K w_k p_k(x) \quad (6.1)$$

$$\iff p(x) = p_z(x), \quad z \sim \text{Cat}(w_1, \dots, w_K). \quad (6.2)$$

In fact $p_z(x)$ is a conditinal distribution, $p(x|z)$, and combined with the distribution of Z we can define the joint

$$p(x, z) := p_z(x)p(z)$$

The the case of a statistical model, data \mathcal{D} is fitted by the mixture model by tuning the model parameters $\theta = \{w, \text{paramers for } p_i\}$. Then the joint distribuion $p(\mathcal{D}, z|\theta)$ is refered as the *complete-data* likelihood in the EM algorithm.

$$p(\mathcal{D}, z|\theta) := p(\mathcal{D}|z, \theta)p(z|\theta)$$

When fitting model parameters we essentially want to find the parameters, that maximize the probability of the parameters given the data, $p(\theta|\mathcal{D})$. Assuming an uninformative/flat prior $p(\theta)$,

$$\begin{aligned} p(\theta|\mathcal{D}) &= \frac{p(\mathcal{D}|\theta)p(\theta)}{p(\mathcal{D})} \\ \Rightarrow \arg \max_{\theta} p(\theta|\mathcal{D}) &= \arg \max_{\theta} p(\mathcal{D}|\theta) \end{aligned}$$

we arrive at the maximum likelihood estimate (MLE). The task of finding the MLE is conviniently done using EM algorithm, since we can look at the likelihood as the marginalized *complete-data* likelihood,

$$p(\mathcal{D}|\theta) = \sum_z p(\mathcal{D}, z|\theta)$$

Background: Expectation-maximization EM <based on [5]>

Expectation maximization is a convinient method for finding ML (or MAP) estimate of a latent variable model. We consider a probabilistic model paramitised with θ ,

$$p(\mathbf{X}, \mathbf{Z}|\theta)$$

where we denote all latent variables \mathbf{Z} , and observed variables \mathbf{X} . Our goal is to find the maximum of the likelihood,

$$p(\mathbf{X}|\theta) = \int p(\mathbf{X}, \mathbf{Z}|\theta) \mu(d\mathbf{Z})$$

maximizing the likelihood itself $p(\mathbf{X}|\theta)$ is assumed difficult but maximizing of the *complete-data* likelihood $p(\mathbf{X}, \mathbf{Z}|\theta)$ is much easier. The algorithm iterates over two steps: The expection (E) step and the maximization (M) step, defined in the following way for iteration t ,

E-step

Define the functional $Q(\theta, \theta^{(t)})$, to be the expected value of the complete-data log likelihood (log likelihood function of θ), with respect to the only random quantaty \mathbf{Z} , which is assumed to follow a distribuion with the density $p(\mathbf{Z}|\mathbf{X}, \theta^{(t)})$, i.e. the conditional distribution of \mathbf{Z} given \mathbf{X} and the current parameter point estimate $\theta^{(t)}$:

$$Q(\theta, \theta^{(t)}) := E_{p(\mathbf{Z}|\mathbf{X}, \theta^{(t)})}[\log p(\mathbf{X}, \mathbf{Z}|\theta)]$$

M-step

After the E-step we find the point estimate $\theta^{(t+1)}$ which maximizes $Q(\cdot|\theta^{(t)})$, i.e.

$$\theta^{(t+1)} = \arg \max_{\theta} Q(\theta|\theta^{(t)})$$

(local) maximization of $p(\mathcal{D}|\theta)$

Input: dataset \mathcal{D} , joint model $p(\mathcal{D}, \mathbf{Z}|\theta)$

while Not converged **do**

$$Q(\cdot, \theta^{(t)}) \leftarrow E_{p(\mathbf{Z}|\mathcal{D}, \theta^{(t)})}[\log p(\mathcal{D}, \mathbf{Z}|\cdot)] \quad \triangleright \text{E-step}$$

$$\theta^{(t+1)} \leftarrow \arg \max_{\theta} Q(\theta|\theta^{(t)}) \quad \triangleright \text{M-step}$$

return: $\theta^{(end)}$

Proof of correctness

We will now give a short proof that maximizing $Q(\cdot|\theta^{(t)})$ maximizes the likelihood $p(\mathbf{X}|\theta)$, where we assume that \mathbf{Z} is a random vector with a discrete distribution. This allow us to use Gibbs inequality:

$$\sum_z p_1(z) \log p_1(z) \geq \sum_z p_1(z) \log p_2(z)$$

where $p_1(\cdot)$ and $p_2(\cdot)$ are densities belonging to two discrete distributions of Z , equality if $p_1(\cdot) = p_2(\cdot)$. From now on we will alter the subscript on the expectations, just have in mind that

$$E_{\theta^{(t)}}[g(Z)] := E_{p(\mathbf{Z}|\mathbf{X}, \theta^{(t)})}[g(Z)] = \sum_z g(z)p(\mathbf{z}|\mathbf{X}, \theta^{(t)})$$

Now to the proof: From bayes rule $p(\mathbf{X}|\theta) = \frac{p(\mathbf{X}, \mathbf{Z})}{p(\mathbf{Z})}$ we can write

$$\log p(\mathbf{X}|\theta) = \log p(\mathbf{X}, \mathbf{Z}) - \log p(\mathbf{Z}|\mathbf{X}, \theta)$$

Now, taking the expectation of the above w.r.t. $p(\mathbf{Z}|\mathbf{X}, \theta^{(t)})$, yields,

$$\begin{aligned} \log p(\mathbf{X}|\theta) &= E_{\theta^{(t)}}[\log p(\mathbf{X}, \mathbf{Z}|\theta)] - E_{\theta^{(t)}}[\log p(\mathbf{Z}|\mathbf{X}, \theta)] \\ &= Q(\theta, \theta^{(t)}) + E_{\theta^{(t)}}[\log p(\mathbf{Z}|\mathbf{X}, \theta)] \end{aligned}$$

Since the above equation holds for any θ , it also holds for $\theta^{(t)}$ now we have,

$$\log p(\mathbf{X}|\theta^{(t)}) = Q(\theta^{(t)}, \theta^{(t)}) + E_{\theta^{(t)}}[\log p(\mathbf{Z}|\mathbf{X}, \theta^{(t)})]$$

Subtracting the two equations, we get,

$$\log p(\mathbf{X}|\theta) - \log p(\mathbf{X}|\theta^{(t)}) = Q(\theta, \theta^{(t)}) - Q(\theta^{(t)}, \theta^{(t)}) + E_{\theta^{(t)}}[\log p(\mathbf{Z}|\mathbf{X}, \theta)] - E_{\theta^{(t)}}[\log p(\mathbf{Z}|\mathbf{X}, \theta^{(t)})]$$

From Gibb's inequality we have that $E_{\theta^{(t)}}[\log p(\mathbf{Z}|\mathbf{X}, \theta^{(t)})] \leq E_{\theta^{(t)}}[\log p(\mathbf{Z}|\mathbf{X}, \theta)]$ where equality only holds for $\theta^{(t)} = \theta$, giving

$$\log p(\mathbf{X}|\theta) - \log p(\mathbf{X}|\theta^{(t)}) \geq Q(\theta, \theta^{(t)}) - Q(\theta^{(t)}, \theta^{(t)})$$

so optimizing $Q(\theta, \theta^{(t)})$ will optimize $\log p(\mathbf{X}|\theta)$ as least as much.

EM for Gaussian mixture

For a Gaussian mixture the p_k distributions in (6.1) is substituted by Gaussian pdfs, i.e. $p_k(x) = \mathcal{N}(x|\mu_k, \Sigma_k)$ and the density of a categorical distribution is $p(z) = \sum_{k=1}^K 1_{z=k} w_k = w_z$, combining the two we get the joint distribution,

$$p(x, z|w, \mu, \Sigma) = w_z \mathcal{N}(x|\mu_z, \Sigma_z)$$

Taking the log and defining $\theta = \{\mu_1, \Sigma_1, w_1, \dots, \mu_K, \Sigma_K, w_K\}$, and assuming iid data

$$\log p(X, Z|\theta) = \sum_i^n (\log(w_{z_i}) + \log(\mathcal{N}(x_i|\mu_{z_i}, \Sigma_{z_i})))$$

Now we are ready to calculate $Q(\theta, \theta^{(t)})$, by taking the expectation of the complete-data log likelihood with respect to the distribution, $p(Z|X, \theta)$,

$$\begin{aligned} E_{p(Z|X, \theta^{(t)})}[\log p(X, Z|\theta)] &= \sum_i^n E_{p(Z|X, \theta^{(t)})}[p(X_i, Z_i|\theta)] \\ &= \sum_i^n E_{p(z_i|x_i, \theta^{(t)})}[p(x_i, z_i|\theta)] \end{aligned}$$

Expectation with respect to unnecessary variables

the last equation holds since taking expectation over a function of a random variable x with respect to a that random variable and more random variables, x, y , is equivalent to the expectation with respect to just x , i.e.

$$\begin{aligned} E_{x,y}[g(x)] &= \int \int g(x) p(x, y) dy dx \\ &= \int g(x) \int p(x, y) dy dx \\ &= \int g(x) p(x) dx = E_x[g(x)] \end{aligned}$$

the posterior distribution is calculated the following way,

$$\begin{aligned} p(z|x, \theta^{(t)}) &= \frac{p(x, z|\theta^{(t)})}{p(x|\theta^{(t)})} \\ &= \frac{p(x, z|\theta^{(t)})}{\sum_z p(x, z|\theta^{(t)})} \\ &= \frac{w_z^{(t)} \mathcal{N}(x|\mu_z^{(t)}, \Sigma_z^{(t)})}{\sum_{k=1}^K w_k^{(t)} \mathcal{N}(x|\mu_k^{(t)}, \Sigma_k^{(t)})} \end{aligned}$$

For simplification we will denote, $\gamma^{(t)}(z_i) := p(z_i|x_i, \theta^{(t)})$, interpreted as the probability of data-point x_i to belong to class z_i . Bishop [5] calls this probability function the *responsibility*. We can now conclude the **E-step**.

$$\begin{aligned}
Q(\cdot, \theta^{(t)}) &= \sum_{i=1}^n p(x_i, z_i | \cdot) \gamma^{(t)}(z_i) \\
&= \sum_{i=1}^n \left[\gamma^{(t)}(z_i) \log(\cdot | z_i) + \gamma^{(t)}(z_i) \log(\mathcal{N}(x_i | \cdot, z_i)) \right]
\end{aligned}$$

or more concretely $\theta = \{\mu_1, \Sigma_1, w_1, \dots, \mu_K, \Sigma_K, w_K\}$,

$$Q(\theta, \theta^{(t)}) = \sum_{i=1}^n \gamma^{(t)}(z_i) \log(w_{z_i}) + \gamma^{(t)}(z_i) \log(\mathcal{N}(x_i | \mu_{z_i}, \Sigma_{z_i})).$$

$Q(\cdot, \theta^{(t)})$ is a concave function - the Gaussian is log-concave and a sum of concave functions is also concave - so it is sufficient and necessary to find its maxima by the root of its derivative,

$$\frac{d}{d\theta} Q(\theta^*, \theta^{(t)}) = 0 \iff \theta^* = \arg \max_{\theta} Q(\theta, \theta^{(t)})$$

Giving the updates...

EM for SPN

6.6 Gaussian Mixture Regression

The Gaussian mixture is a generative model of the joint probability of x and y given as,

$$p(x, y) = \sum_{k=1}^K \pi^{(k)} \mathcal{N}(x, y | \mu^{(k)}, \Sigma^{(k)}), \quad \mu = \begin{bmatrix} \mu_x \\ \mu_y \end{bmatrix}, \quad \Sigma = \begin{bmatrix} \Sigma_{xx} & \Sigma_{xy} \\ \Sigma_{yx} & \Sigma_{yy} \end{bmatrix}$$

This is trained using the EM algorithm. We will now show how the conditinal is calculated exact.

6.6.1 GMR - prediction

We need the conditional distribution of the Gaussian mixture in order to get the predictive distribution. We will now formulate the conditinal distribution in terms of conditional and marginals of the individual mixture components. First of all the marginal distribution $p(x)$ of the mixture is given as, **How??!**

$$p(x) = \sum_{k=1}^K \pi^{(k)} \mathcal{N}(x | \mu_x^{(k)}, \Sigma_{xx}^{(k)})$$

next the joint distribution can be decomposed with the probability chain rule,

$$\begin{aligned}
p(x, y) &= p(x) p(y|x) \\
\implies \mathcal{N}(x, y | \mu, \Sigma) &= \mathcal{N}(x | \mu_x, \Sigma_{xx}) \mathcal{N}(y | \mu_{y|x}, \Sigma_{y|x})
\end{aligned}$$

And we can formulate the conditinal in terms of individual multivariate Gaussians,

$$p(y|x) = \frac{p(y, x)}{p(x)} \tag{6.3}$$

$$= \sum_{k=1}^K \frac{\pi^{(k)}}{p(x)} \mathcal{N}(x, y | \mu^{(k)}, \Sigma^{(k)}) \tag{6.4}$$

$$= \sum_{k=1}^K \frac{\pi^{(k)} \mathcal{N}(x | \mu_x^{(k)}, \Sigma_{xx}^{(k)})}{p(x)} \mathcal{N}(y | \mu_{y|x}^{(k)}, \Sigma_{y|x}^{(k)}) \tag{6.5}$$

$$= \sum_{k=1}^K \pi_{y|x}^{(k)} p(y|x, \mu_{y|x}^{(k)}, \Sigma_{y|x}^{(k)}) \tag{6.6}$$

where $\pi_{y|x}^{(k)} := \frac{\pi^{(k)} \mathcal{N}(x|\mu_x^{(k)}, \Sigma_{xx}^{(k)})}{\sum_{i=1}^K \mathcal{N}(x|\mu_x^{(i)}, \Sigma_{xx}^{(i)})}$. So we see that the conditional of a Gaussian mixture model is again a Gaussian mixture model.

Background: Conditional of multivariate Gaussian

The conditional distribution is defined as [5] ..?

$$p(y|x, \mu, \Sigma) = \mathcal{N}(y|\mu_{y|x}, \Sigma_{y|x}) \quad \mu = \begin{bmatrix} \mu_x \\ \mu_y \end{bmatrix}, \quad \Sigma = \begin{bmatrix} \Sigma_{xx} & \Sigma_{xy} \\ \Sigma_{yx} & \Sigma_{yy} \end{bmatrix}$$

where

$$\mu_{y|x} = \mu_y + \Sigma_{yx} \Sigma_{xx}^{-1} (x - \mu_x) \tag{6.7}$$

$$\Sigma_{y|x} = \Sigma_{yy} - \Sigma_{yx} \Sigma_{xx}^{-1} \Sigma_{xy} \tag{6.8}$$

6.6.2 GMR - Learning

EM-algorithm, classic example.

7 Results

<reproducer resultater fra Arayns thesis> Egne forsøg med SPN og mixture regression

7.0.1 regression benchmark

We will benchmark our different regression models against a simple empirical mean and empirical std. normal distribution benchmark,

$$p(y|x\mathcal{D}) = \mathcal{N}(y|\bar{\mathbf{y}}, \bar{\sigma}^2(\mathbf{y}))$$

where $\bar{\mathbf{y}} = \frac{1}{n} \sum_{i=1}^n \mathbf{y}_i$ and $\bar{\sigma}^2(\mathbf{y}) = \frac{1}{n-1} \sum_{i=1}^n (\mathbf{y}_i - \bar{\mathbf{y}})^2$ this is however not a good model for Bayesian optimization, as it will not provide a new candidate point, as all points are equally good.

7.1 Mixture regression on simple functions

Choosing a good set of design parameters is crucial, the manipulation...

Figure 7.1: Regression plot for one problem with all dims: Number of data points / dims, mean squared error, with bayesline mean prediction

Figure 7.2: Regression plot for another problem

Problem 2D	Winner (rel error)	Second	Third
f1	GP (0.2)	Neural Network (0.00001)	Naive (0.002)
f2	Naive (0.002)		
f3	Neural Network (0.00001)		
f4			

Figure 7.3: Table summarizing for all problems

Name	Age	Nickname
MrXavierHuon	32	Xav'
MrBaptisteClement	1	Baby
MmeLouiseBourgeau	28	Lou Loue

Table 7.1: An example table.

Figure 7.4: BayesOpt plot: Number of iterations, distance to optima using EI, with bayesline random search

Bibliography

- [1] Mashall Aryan. “Sample-efficient Optimization Using Neural Networks”. <https://researcharchive.vuw.ac.nz/handle/10063/9035>. PhD thesis. Victoria University of Wellington, New Zealand, 2020.
- [2] Jasper Snoek et al. *Scalable Bayesian Optimization Using Deep Neural Networks*. 2015. arXiv: 1502.05700 [stat.ML].
- [3] Jost Tobias Springenberg et al. “Bayesian Optimization with Robust Bayesian Neural Networks”. In: *Advances in Neural Information Processing Systems*. Ed. by D. Lee et al. Vol. 29. Curran Associates, Inc., 2016. url: <https://proceedings.neurips.cc/paper/2016/file/a96d3afec184766bfeca7a9f989fc7e7-Paper.pdf>.
- [4] Roman Garnett. *Bayesian Optimization*. in preparation. Cambridge University Press, 2022.
- [5] Christopher M Bishop and Nasser M Nasrabadi. *Pattern recognition and machine learning*. Vol. 4. 4. Springer, 2006.

A Implementation

Bayesian optimization is an effective framework for finding optimizers of a highly expensive (in terms of money, time, human attention or computational processing) or noisy objective function. First a prior is defined over possible functions and updated to a posterior according to already obtained observations/samples of the objective function. Next an acquisition function uses this posterior, also called a surrogate model, and is then utilized to find the next location in the optimization landscape to sample from. The far most common surrogate model is Gaussian Process (GP), partially due to its ability to represent its posterior in closed form. However, it also comes with shortcomings: Its inference, although it is exact, scales cubic with amount of samples and it imposes strong assumptions of a well behaved objective function.

This thesis aims to investigate surrogate models different from GPs in order to improve on either the accuracy of the surrogate model or the inference cost of it. Meanwhile Bayesian Neural Networks (BNN) already have proven useful as surrogate models [1][2][3] (with cost of inference, which scales linearly and less strong assumptions) this thesis additionally wants to investigate Sum Product networks (SPN). An SPN is - similarly to a BNN - a deep probabilistic model and still expressive but with tractable inference, which potentially could lead to advantages over BNNs.

Technical
University of
Denmark

Building 321
2800 Kgs. Lyngby
Tlf. 4525 1700

www.github.com/SimonKruse0/master-thesis